



# Parallelizing the Web Browser

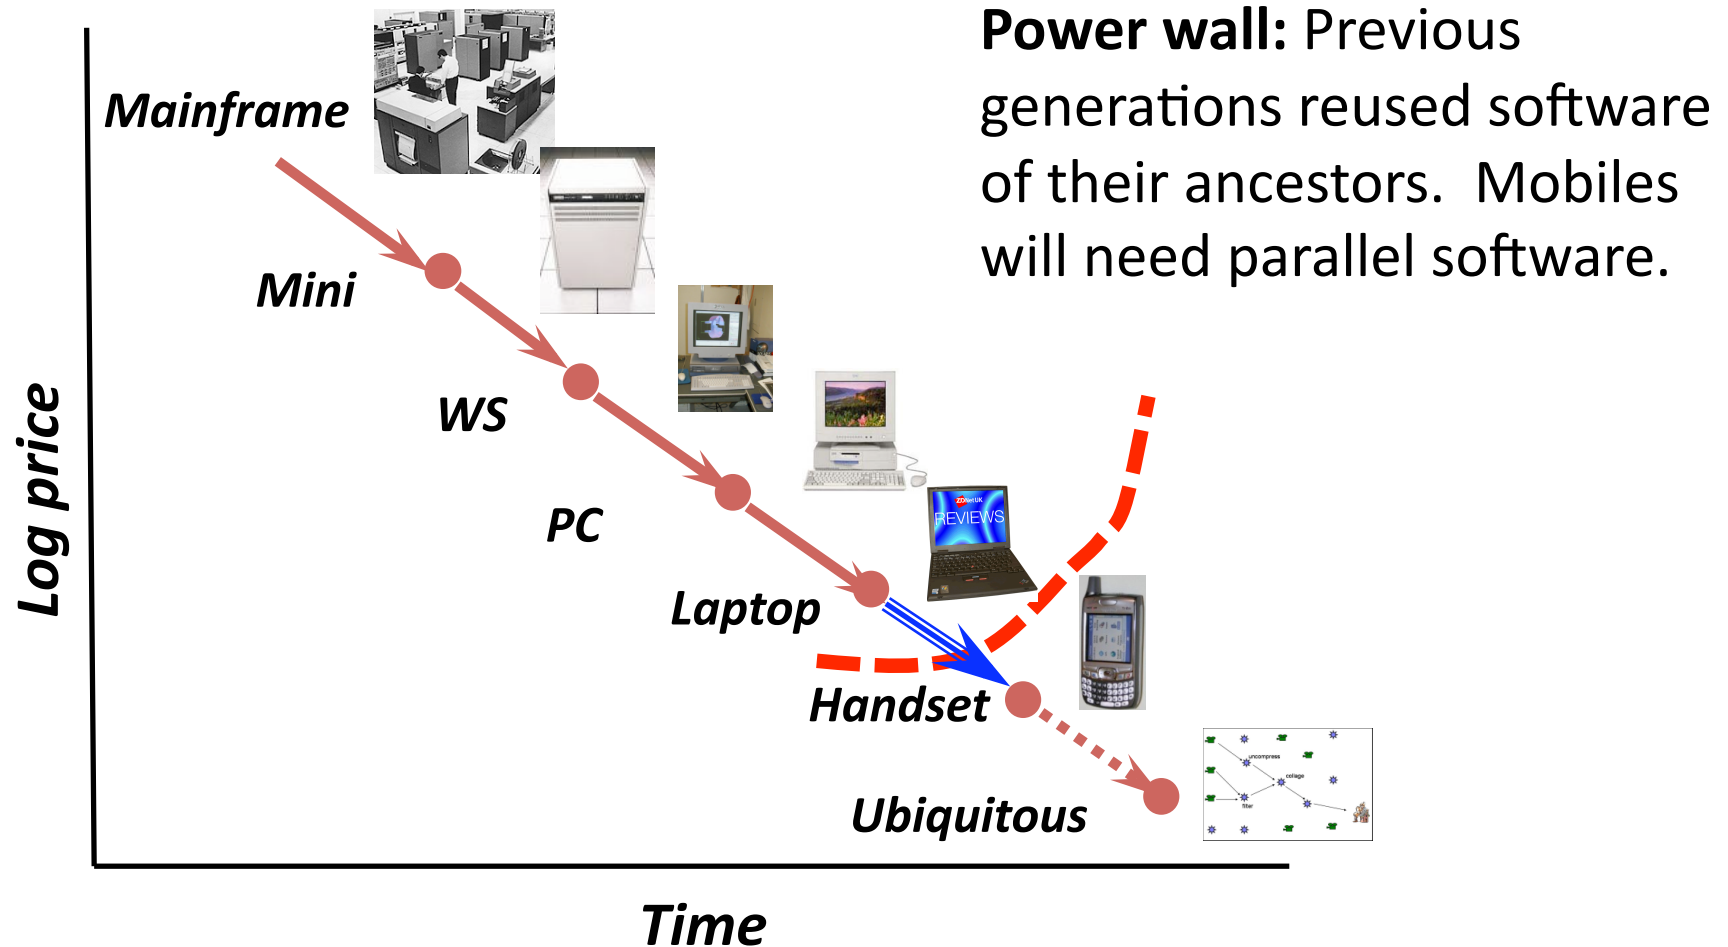
Chris Jones, Rose Liu, [Leo Meyerovich](#)  
Krste Asanovic, and Rastislav Bodik

ParLab  
UC Berkeley






# The Transition to Handhelds



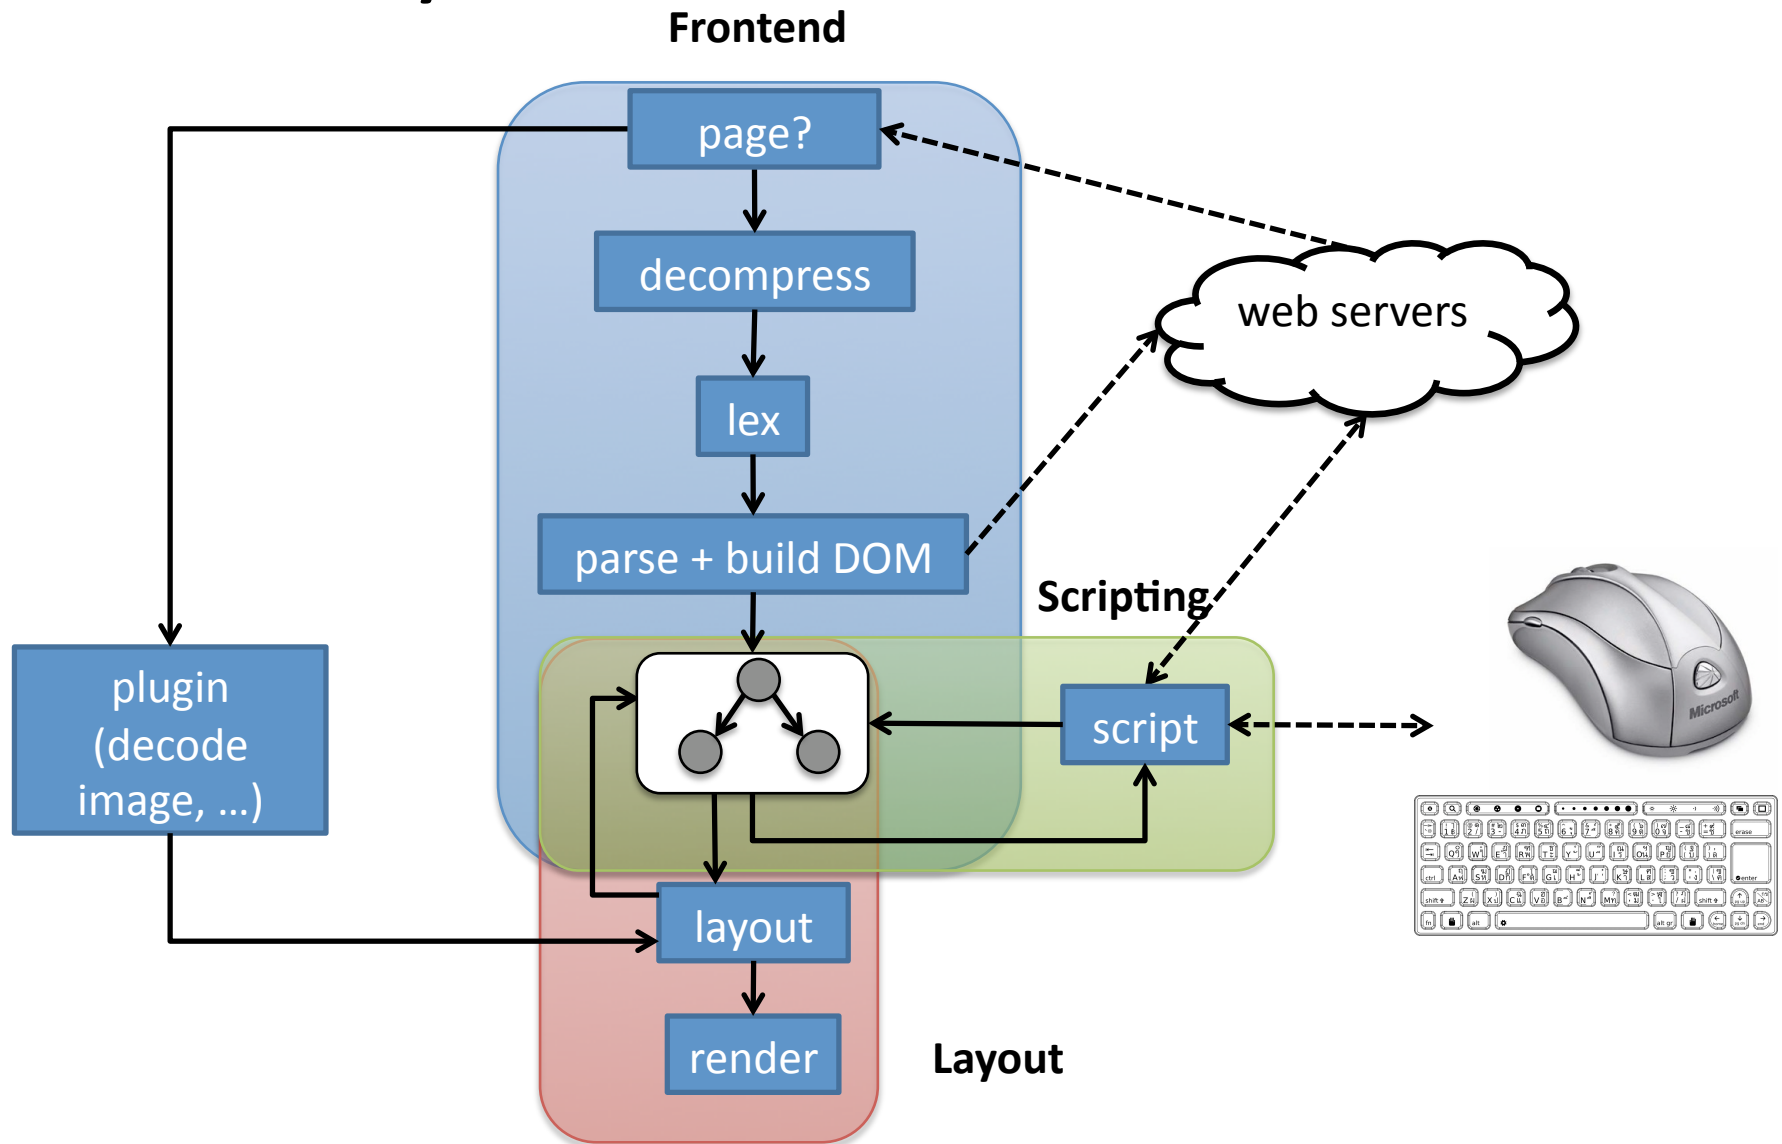

Soon on mobile: 4-cores x 2-threads x 8-SIMD = 64-way parallelism


# Why Parallelize a Browser?



- **Dominant application platform**
  - easy deployment: apps downloaded, JS portable
  - productive programming: scripting, layout
- **... but not on handhelds**
  - native frameworks for: iPhone, Google Android
  - slow: for Slashdot, Laptop: 3s => iPhone: 21s
- **Parallel browser may need new architecture**
  - ex: JavaScript relies on “gotos”, is too serial

# Anatomy of a Browser





# Project Status

## 1. Developed *work-efficient* algorithms

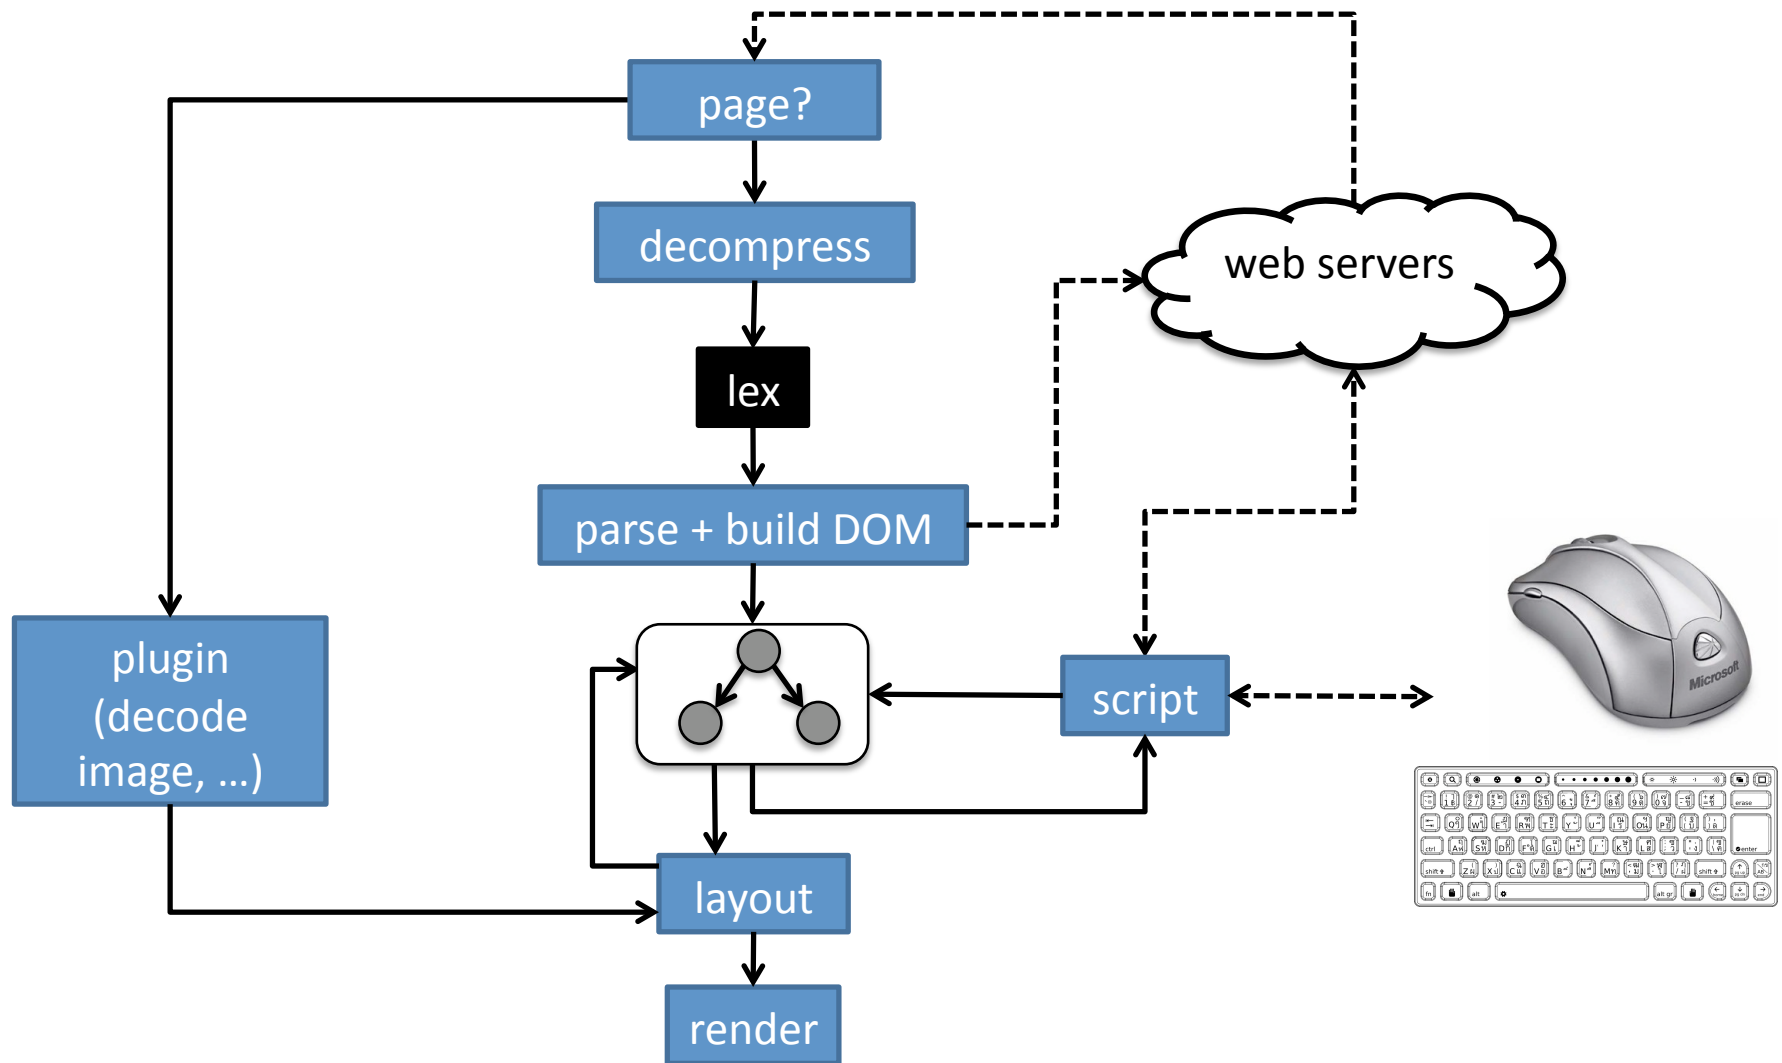

work-efficient : no more work than sequential algo.

- *layout*: parallel-map with a tiling optimization
- *layout*: break up tree traversal into five parallel ones
- *lexing*: speculation to break sequential dependencies


## 2. Reexamining the scripting programming model

- *programmer productivity*: from callbacks to actors
- *performance*: adding structure to detect dependences

# Frontend: Lexing




# Lexing, from 10,000 feet




**Goal:** given lexical spec and input, find lexemes

STag ::= < [ ^ > ] \* >  
 Content ::= [ ^ < ] +  
 ETag ::= < / [ ^ > ] \* >




*STag*

<	b	>	B	e	r	k	e	l	e	y	!	<	/	b	>
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---




(label each character with its state)


# Inherently Sequential?



STag ::=  $\langle [\wedge >]^* \rangle$   
 Content ::=  $[\wedge <]^+$   
 ETag ::=  $\langle / [\wedge >]^* \rangle$








# An observation

In lexing, irrespective of where DFA starts, it converges to a *stable, recurring* state


Lexing:

<	b	>	B	e	r	k	e	l	e	y	!	<	/	b	>
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---




Parallel scans thus need not scan from all possible states, just one, yielding a work-efficient algorithm.


# Our solution (1/2): Partition




- split input into blocks with  $k$ -character overlap
- scan in parallel; start block from a *tolerant* state




# Our solution (2/2): Speculate




- split input into blocks with  $k$ -character overlap
- scan in parallel; start block from a *tolerant* state
- check if blocks converge: expected in  $k$ -overlap
- speculation may fail; if so, block is rescanned





# Speedup: Flex vs Cell

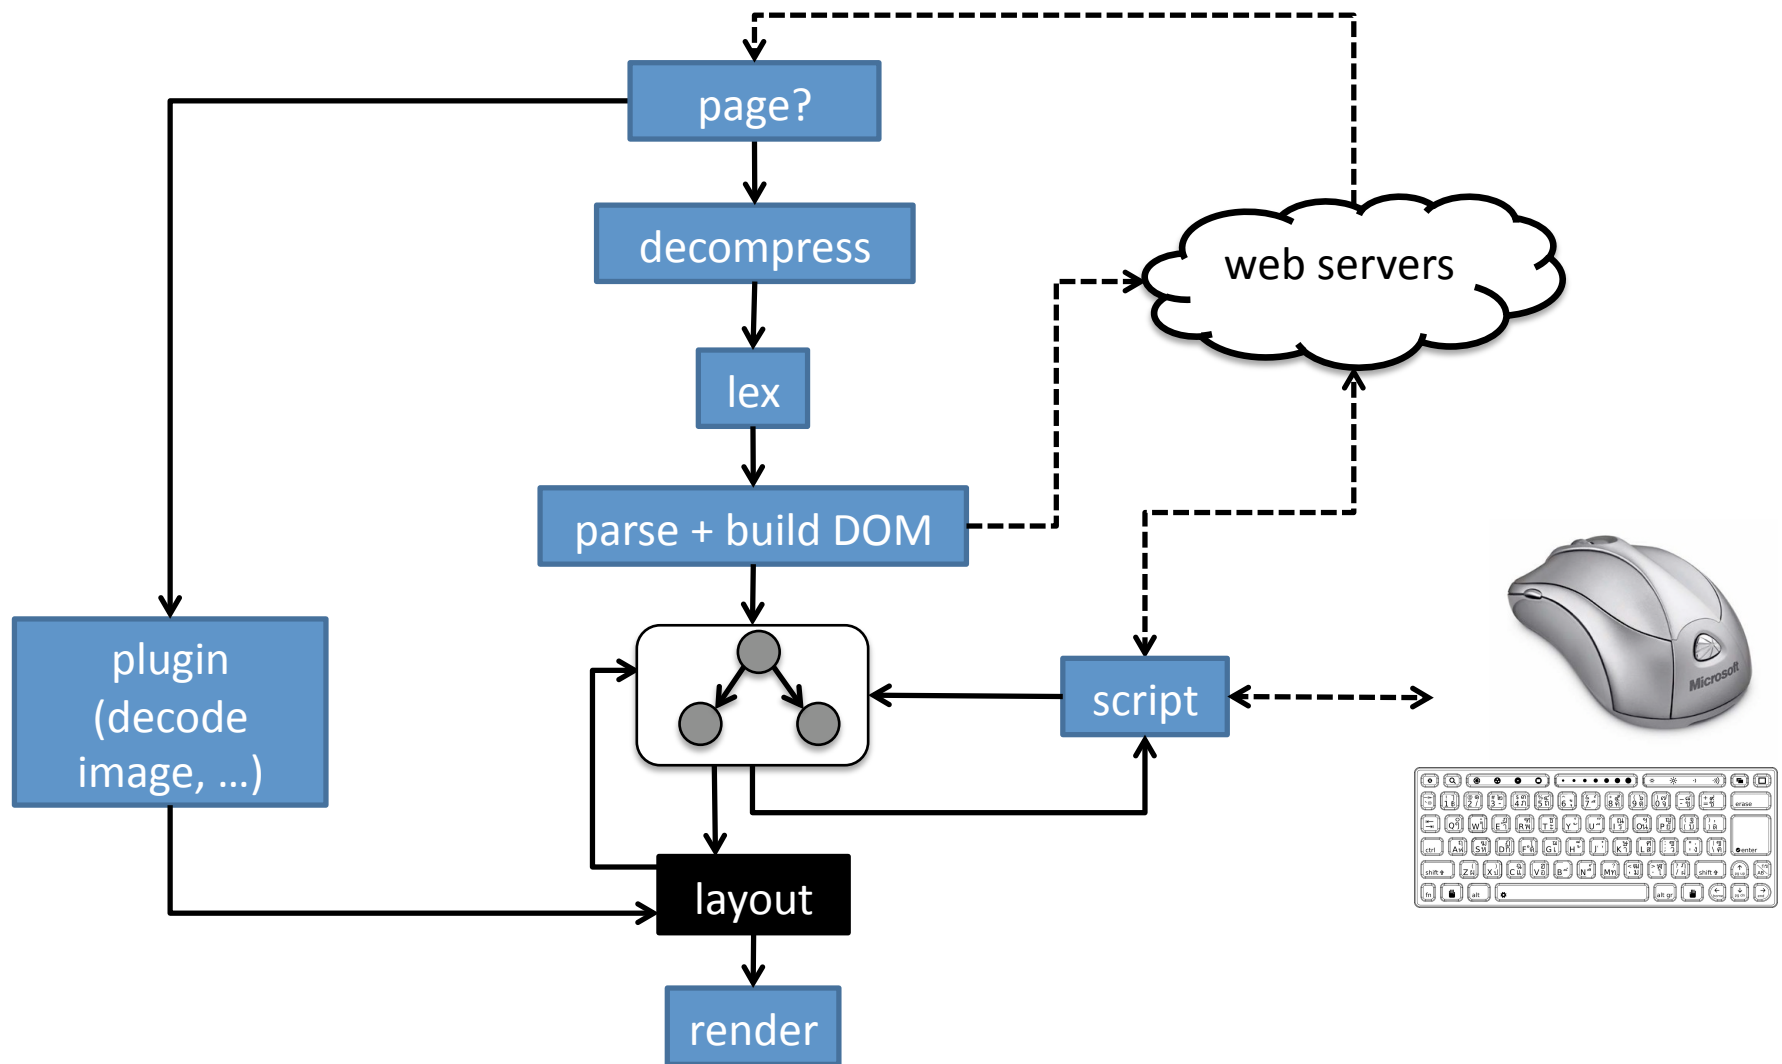

Speedup over flex for various numbers of cores




*today's page  
sizes: 5 cores  
are 4.5x faster  
than flex*

**baseline:** (sequential) flex on the CELL main CPU

# Layout Solving (1/2)




# Rule Matching



**Goal:** Match rules with nodes:


- a rule: p img { fontsize: 7px }
- match tag path
- path-rule matching
  - end with the same node
  - and are a substring



<b>selectors</b>	p	img	p img
<b>properties</b>	height=83%	width=100px float=left	fontsize=7px

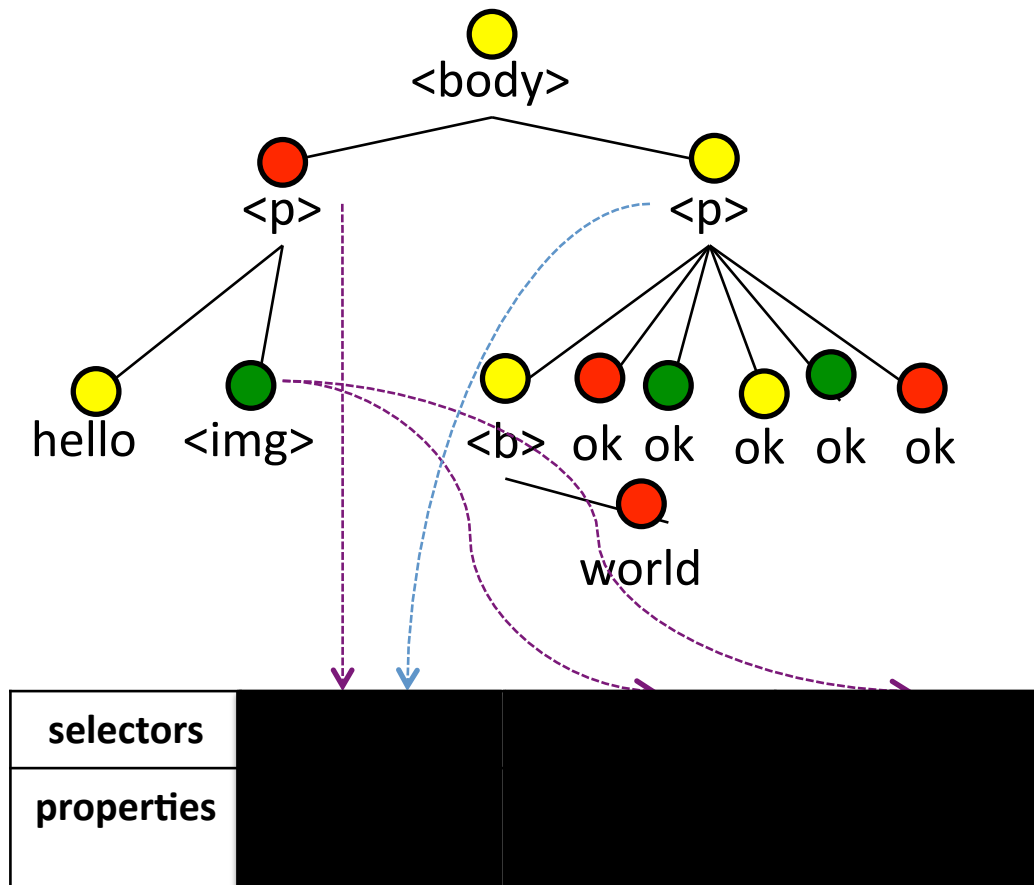
- 

<b>selectors</b>	p	img	p img
<b>properties</b>	height=83%	width=100px float=left	fontsize=7px




# Tiling for Caches

**Problem: all the nodes + selectors might not fit in cache!**






# Speedup (Cilk++)





## Speedup vs. Fastest Sequential (Slashdot)

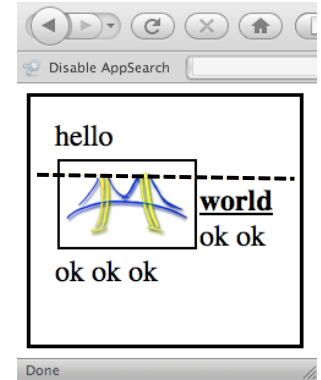




2 socket x 4 core x 2 thread (2.6 Ghz, 12x 1 GB)

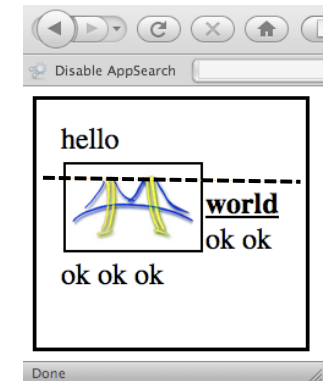


# Layout Solving (2/2)





# Problem: Layout a Page




# It looks rather sequential..



# But not entirely



# 5 Phases: Each Exhibits Tree Parallelism



**Phase 1: font size, temporary width**



**Phase 2: preferred max & min width**

**Phase 3: solved width**

**Phase 4: height, relative x/y position**

**Phase 5: absolute x/y position**

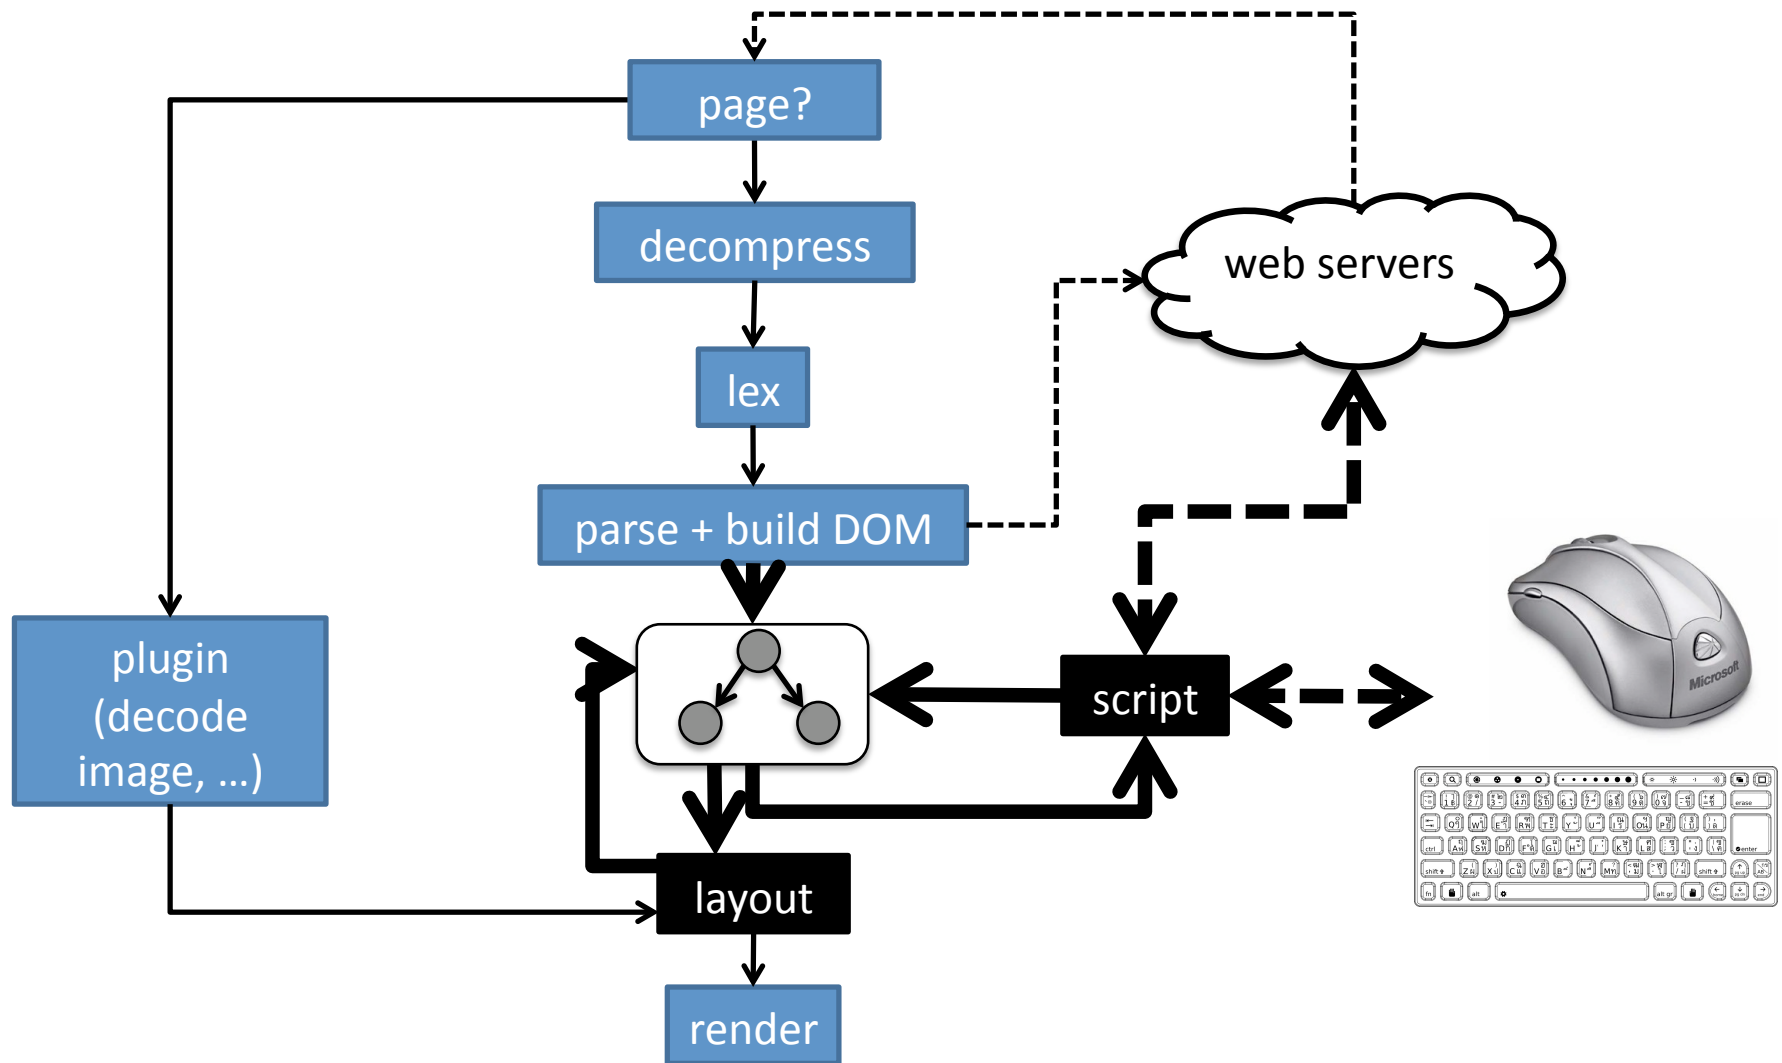

# Results: layout (modeled)



**Baseline:** Cilk++ model on 1 core.

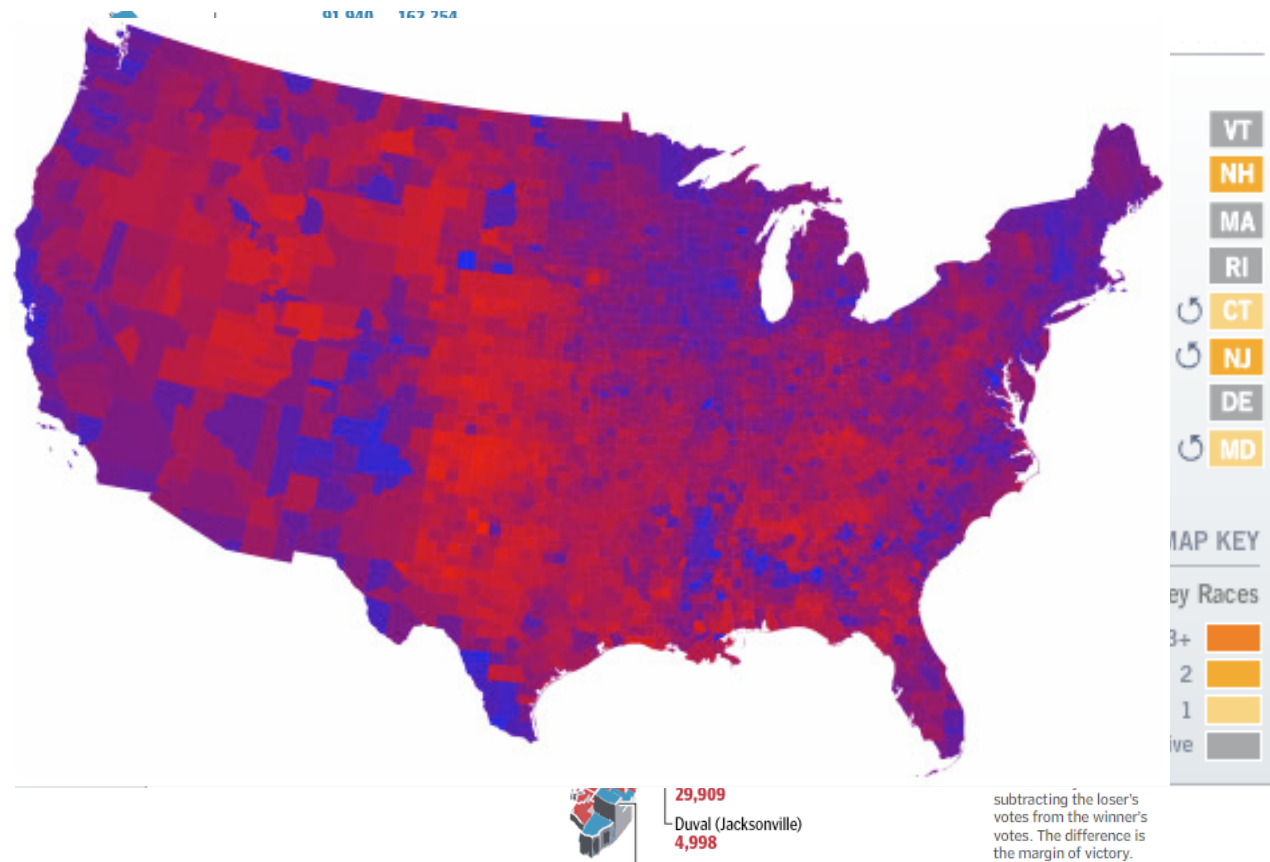
- Eight socket x 4 core AMD Opteron 2356 Barcelona Sun X4600
- Dual socket x 4 core AMD Opteron 2356 Barcelona Sun X2200
- Preproduction 2 socket x 4 core x 2 thread Intel Xeon Nehalem

# Scripting






# Why parallelize scripting (example)



**Example:** animate between different views

- each transition: recolor, resize each state or county
- animation rate 30fps => 33ms for 1000s of nodes

# The browser programming model



- Nonpreemptive event model
- Handlers respond to events
- Handlers execute atomically
  - document changes cause relayout
  - style changes cause relayout
- To parallelize, must understand how the document is shared
  - document-carried dependencies:  
handler A: `california.x = 100;`  
handler B: `var z = california.x;`
  - layout-carried dependencies:  
handler A: `america.w = 200%;`  
layout: `california.w = 200%;`  
handler B: `var z = california.w;`

# Concurrency bugs

1. GUI animations and interactions
  - several animations modifying an object simultaneously
2. Server interactions
  - responses to requests may be delayed, reordered
3. Eager script loading
  - executing a script on a document before done loading

# “Gotos” in JavaScript


```
<div id="box" style="position:absolute; background: yellow;">  
  My box  
</div>
```


```
<script>  
document.addEventListener (  
  'mousemove',  
  function (e) {  
    var left = e.pageX;  
    var top = e.pageY;  
    setTimeout(function() {  
      document.getElementById("box").style.top = top;  
      document.getElementById("box").style.left = left;  
    } , 500);  
  }, false);  
</script>
```

# Preliminary design of our language

Program structure is clearer when data and control is explicit

- in dataflow version: **changing mouse coordinates are streams**
- coordinate streams adjust box position after they are delayed
- **structured names** of document element allow analysis





# Summary

## 1. Developed *work-efficient* algorithms

- *Rule matching*: parallel-map with a tiling optimization
- *Layout*: break up tree traversal into five parallel ones
- *Lexing*: speculation to break sequential dependencies

## 2. Reexamining the scripting programming model

- *programmer productivity*: from callbacks to actors
  - influenced by Flapjax, Ptolemy, Max/MSP, LabVIEW
- *performance*: adding structure to detect dependences
  - current browsers: JIT compilation, font vectorization, task parallelism eg for image rendering – all these are useful, too.